

Overview: Our project application is *The Consult* - an AI-powered medical literature search application tailored for clinicians and researchers. We have data collection, a ChromaDB database to support RAG, a fine-tuned LLM, a frontend UI, APIs, and application deployment to support this application. This is the code organization:

```
src/
  ├── datapipeline/ # Data collection directory
  ├── finetuning/ # Model fine-tuning directory
  ├── frontend/ # Frontend UI directory
  ├── llm-api/ # Application API directory
  ├── models/ # RAG module directory
  └── secrets/ # GCP credentials directory
```

Data Collection: We collect data for RAG as follows:

- 1) Download and parse Pubmed XML files containing metadata and abstracts
- 2) Upload the pickled abstracts and metadata to a local directory or a GCS bucket

Technologies Used: Python, Docker, GCS

ChromaDB Database: We insert the data from the data collection step into a ChromaDB collection for RAG as follows:

- 1) Read data from GCS and create semantic chunks and embeddings from the abstracts
- 2) Upload chunks, embeddings, and associated metadata to a ChromaDB collection

Technologies Used: Python, Docker, GCS, VertexAI for embeddings, DVC for data versioning

Fine-tuned LLM: We fine-tuned a Gemini 2.5 model to customize responses towards clinicians or researchers. This involves:

- 1) Generating a question-answer dataset and converting it to the appropriate format
- 2) Fine-tuning Gemini 2.5 using the generated dataset

Technologies Used: Python, Docker, VertexAI for Gemini and finetuning

Frontend UI: Our frontend displays a search bar to ask medical questions, metadata filtering options, and an LLM response with references. The user can toggle between Clinical Practice and Research mode.

Technologies Used: React, Typescript, Vite, CSS

APIs: We use a FastAPI web service to connect the backend (ChromaDB database and Google Gemini) to our frontend UI. This is how the APIs work:

- 1) Accept questions and metadata from users in the frontend and send results to API server
- 2) Server sends user's request to the RAG module and outputs an LLM response
- 3) Return LLM response and citations back to the frontend

Technologies Used: Python (FastAPI), Docker, VertexAI for Gemini

Deployment: We deploy the application that connects ChromaDB, Google Gemini, the APIs, and the frontend together. The application can be deployed locally or on a virtual machine.

Technologies Used: Uvicorn, Terraform, Google Cloud Run, Google Virtual Machine